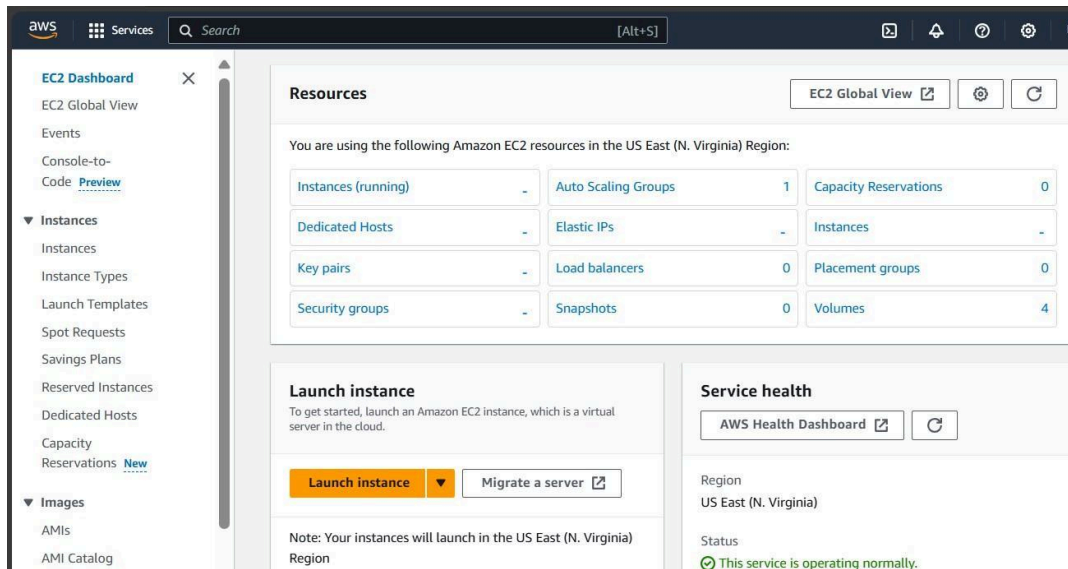


EXPERIMENT NO. 3

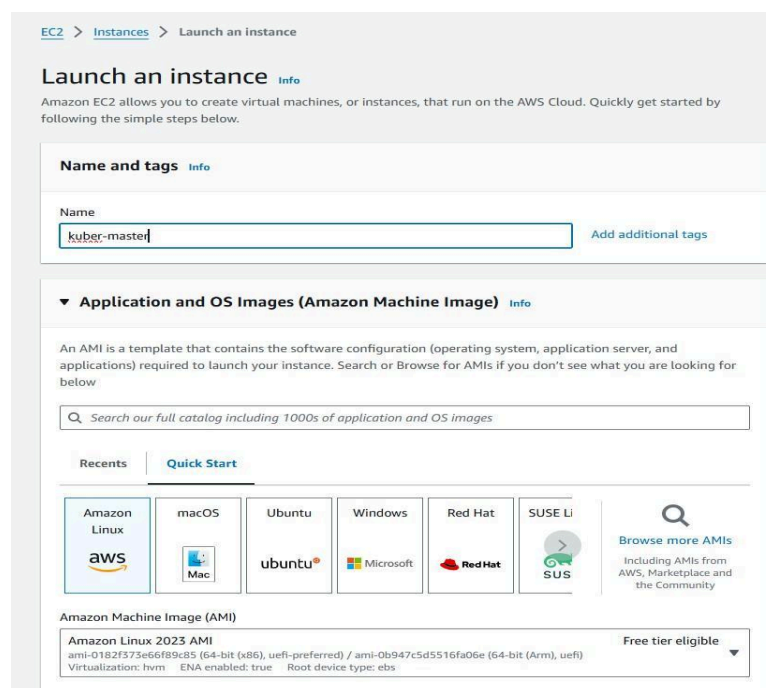
Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud.

Procedure:

1. Creation Of Instance



Search EC-2 instance. Then create three EC-2 instances and choose Amazon Linux as OS and also allow ssh traffic from anywhere.



Description

Amazon Linux 2023 is a modern, general purpose Linux-based OS that comes with 5 years of long term support. It is optimized for AWS and designed to provide a secure, stable and high-performance execution environment to develop and run your cloud applications.

Architecture

64-bit (x86)

Boot mode

uefi-preferred

AMI ID

ami-0182f373e66f89c85

Verified provider

▼ Instance type

Info | Get advice

Instance type

t3.medium
Family: t3 2 vCPU 4 GiB Memory Current generation: true
On-Demand SUSE base pricing: 0.0979 USD per Hour
On-Demand Windows base pricing: 0.06 USD per Hour
On-Demand Linux base pricing: 0.0416 USD per Hour
On-Demand RHEL base pricing: 0.0704 USD per Hour

All generations

Compare instance types

Additional costs apply for AMIs with pre-installed software

▼ Key pair (login)

Info

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

server

Create new key pair

▼ Network settings

Info

Edit

Network

Info

vpc-051bba342b5626898

Subnet

Info

No preference (Default subnet in any availability zone)

Auto-assign public IP

Info

Enable

Additional charges apply when outside of free tier allowance

Firewall (security groups)

Info

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

Create security group

Select existing security group

We'll create a new security group called **launch-wizard-31** with the following rules:

Allow SSH traffic from

Helps you connect to your instance.

Anywhere

0.0.0.0/0

Allow HTTPS traffic from the internet

To set up an endpoint, for example when creating a web server

Allow HTTP traffic from the internet

To set up an endpoint, for example when creating a web server

Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

▼ Configure storage

Info

Advanced

1x

8

GiB

gp3

Root volume (Not encrypted)

Free tier eligible customers can get up to 30 GiB of EBS General Purpose (SSD) or Magnetic storage

▼ Summary

Number of instances

Info

1

Software Image (AMI)

Amazon Linux 2023 AMI 2023.5.2...read more

ami-0182f373e66f89c85

Virtual server type (instance type)

t3.medium

Firewall (security group)

New security group

Storage (volumes)











1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GiB of snapshots, and 100 GB of bandwidth to the internet.

Cancel

Launch instance

Review commands

<input type="checkbox"/>	Name 	Instance ID	Instance state
<input type="checkbox"/>	kube-master	i-00aa79ac09d7462c0	 Running  
<input type="checkbox"/>	kube-worker1	i-0bab86cd3fbfcb40a	 Running  
<input type="checkbox"/>	kube-worker2	i-00dcfd302ffd80dda	 Running  

- ```
Sadneya@DESKTOP-IEPNL3D MINGW64 ~ (master)
$ cd downloads

Sadneya@DESKTOP-IEPNL3D MINGW64 ~/downloads (master)
$ chmod 400 "server.pem"

Sadneya@DESKTOP-IEPNL3D MINGW64 ~/downloads (master)
$ ssh -i "server.pem" ec2-user@ec2-54-174-206-93.compute-1.amazonaws.com
The authenticity of host 'ec2-54-174-206-93.compute-1.amazonaws.com (54.174.206.93)' can't be established.
ED25519 key fingerprint is SHA256:T+tsGyI15gAvUvjeAZ7GjDIWXHOaI4EPF5g5oICrkoQ.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-54-174-206-93.compute-1.amazonaws.com' (ED25519)
to the list of known hosts.
```
- 
- ```
      #_
     ~\ #####~
    ~~ \#####\
       ~\####|
       ~\##/\
        ~\#/_~>
         ~~~
          ~..
           _/_
            /m/'
```
- Amazon Linux 2023
<https://aws.amazon.com/linux/amazon-linux-2023>

2. Installation Of Docker on three machines

- For installation of Docker into all three machines run the following command:

`sudo yum install docker -y`

```
[ec2-user@ip-172-31-87-114 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:06:20 ago on Fri Sep 13 03:20:22 2024.
Dependencies resolved.
=====
Package                        Arch      Version                               Repository      Size
=====
Installing:
docker                        x86_64    25.0.6-1.amzn2023.0.2               amazonlinux     44 M
Installing dependencies:
containerd                    x86_64    1.7.20-1.amzn2023.0.1               amazonlinux     35 M
iptables-libs                 x86_64    1.8.8-3.amzn2023.0.2               amazonlinux    401 k
iptables-nft                  x86_64    1.8.8-3.amzn2023.0.2               amazonlinux    183 k
libcgroup                     x86_64    3.0-1.amzn2023.0.1                 amazonlinux     75 k
libnetfilter_conntrack        x86_64    1.0.8-2.amzn2023.0.2               amazonlinux     58 k
libnfnetlink                  x86_64    1.0.1-19.amzn2023.0.2              amazonlinux     30 k
libnftnl                      x86_64    1.2.2-2.amzn2023.0.2               amazonlinux     84 k
pigz                          x86_64    2.5-1.amzn2023.0.3                 amazonlinux     83 k
runc                          x86_64    1.1.13-1.amzn2023.0.1              amazonlinux    3.2 M

Transaction Summary
=====
Install 10 Packages

Total download size: 84 M
Installed size: 317 M
Downloading Packages:
(1/10): iptables-libs-1.8.8-3.amzn2023.0.2.x86_ 3.6 MB/s | 401 kB    00:00
(2/10): iptables-nft-1.8.8-3.amzn2023.0.2.x86_6 4.6 MB/s | 183 kB    00:00

Installed:
containerd-1.7.20-1.amzn2023.0.1.x86_64
docker-25.0.6-1.amzn2023.0.2.x86_64
iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
libcgroup-3.0-1.amzn2023.0.1.x86_64
libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
libnfnetlink-1.0.1-19.amzn2023.0.2.x86_64
libnftnl-1.2.2-2.amzn2023.0.2.x86_64
pigz-2.5-1.amzn2023.0.3.x86_64
runc-1.1.13-1.amzn2023.0.1.x86_64

Complete!
```

- Then, configure cgroup in a daemon.json file by using following commands
`cd /etc/docker`
`cat <<EOF | sudo tee /etc/docker/daemon.json`

```
{
"exec-opts":
["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
}
```

```
"storage-driver": "overlay2"
}
EOF
```

```
[ec2-user@ip-172-31-87-114 ~]$ cd /etc/docker
[ec2-user@ip-172-31-87-114 docker]$ cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
"storage-driver": "overlay2"
}
EOF
{
"exec-opts": ["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
"storage-driver": "overlay2"
}
}
```

- Then after this run the following command to enable and start docker and also to load the daemon.json file.

```
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
```

```
[ec2-user@ip-172-31-80-126 docker]$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
```

- Then check the version of docker installed. docker -v

```
[ec2-user@ip-172-31-80-126 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc
```

3. Installation Of Kubernetes on three machines

- SELinux needs to be disabled before configuring kubelet thus run the following command
`sudo setenforce 0`

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

```
[ec2-user@ip-172-31-80-126 docker]$ sudo setenforce 0  
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

- Here We are adding kubernetes using the repository whose command is given below.
`cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo`

```
[kubernetes]
```

```
name=Kubernetes
```

```
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
```

```
enabled=1
```

```
gpgcheck=1
```

```
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
```

```
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
```

```
EOF
```

```
[ec2-user@ip-172-31-80-126 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo  
[kubernetes]  
name=Kubernetes  
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/  
enabled=1  
gpgcheck=1  
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key  
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni  
EOF  
[kubernetes]  
name=Kubernetes  
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/  
enabled=1  
gpgcheck=1  
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key  
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
```

- After that Run following command to make the update and also to install kubelet, kubeadm, kubectl: `sudo yum update`

```
[ec2-user@ip-172-31-80-126 docker]$ sudo yum update  
Kubernetes  
Dependencies resolved.  
Nothing to do.  
Complete!
```

sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes

```
[ec2-user@ip-172-31-80-126 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Last metadata expiration check: 0:00:10 ago on Fri Sep 13 10:31:17 2024.
Dependencies resolved.
```

Package	Architecture	Version	Repository	Size
Installing:				
kubeadm	x86_64	1.30.5-150500.1.1	kubernetes	10
kubectl	x86_64	1.30.5-150500.1.1	kubernetes	10
kubelet	x86_64	1.30.5-150500.1.1	kubernetes	17 M
Installing dependencies:				
conntrack-tools	x86_64	1.4.6-2.amzn2023.0.2	amazonlinux	208 k
cri-tools	x86_64	1.30.1-150500.1.1	kubernetes	8.6 M
kubernetes-cni	x86_64	1.4.0-150500.1.1	kubernetes	6.7 M
libnetfilter_cthelper	x86_64	1.0.0-21.amzn2023.0.2	amazonlinux	24 k
libnetfilter_cttimeout	x86_64	1.0.0-19.amzn2023.0.2	amazonlinux	24 k
libnetfilter_queue	x86_64	1.0.5-2.amzn2023.0.2	amazonlinux	30 k
Transaction Summary				
Install 9 Packages				

```
Total
Kubernetes
Importing GPG key 0x9A296436:
  Userid : "Isv:kubernetes OBS Project <isv:kubernetes@build.opensuse.org>"
  Fingerprint: DE15 B144 86CD 3778 9E87 6E1A 2346 54DA 9A29 6436
  From : https://pkgs.k8s.io/core/stable/v1.30/rpm/repodata/repomd.xml.key
Key imported successfully
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
  Preparing :
  Installing : kubernetes-cni-1.4.0-150500.1.1.x86_64
  Installing : cri-tools-1.30.1-150500.1.1.x86_64
  Installing : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64
  Installing : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
  Installing : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
  Installing : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
  Running scriptlet: conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
  Installing : kubelet-1.30.5-150500.1.1.x86_64
  Running scriptlet: kubelet-1.30.5-150500.1.1.x86_64
  Installing : kubeadm-1.30.5-150500.1.1.x86_64
  Installing : kubectl-1.30.5-150500.1.1.x86_64
  Running scriptlet: kubectl-1.30.5-150500.1.1.x86_64
  Verifying : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
  Verifying : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
  Verifying : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
  Verifying : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64
  Verifying : cri-tools-1.30.1-150500.1.1.x86_64
  Verifying : kubeadm-1.30.5-150500.1.1.x86_64
  Verifying : kubectl-1.30.5-150500.1.1.x86_64
  Verifying : kubelet-1.30.5-150500.1.1.x86_64
  Verifying : kubernetes-cni-1.4.0-150500.1.1.x86_64
Installed:
  conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      cri-tools-1.30.1-150500.1.1.x86_64      kubeadm-1.30.5-150500.1.1.x86_64
  kubectl-1.30.5-150500.1.1.x86_64                kubelet-1.30.5-150500.1.1.x86_64      kubernetes-cni-1.4.0-150500.1.1.x86_64
  libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64  libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64  libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64
Complete!
```

- After installing Kubernetes, we need to configure internet options to allow bridging.
 - sudo swapoff -a
 - echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
 - sudo sysctl -p

```
[ec2-user@ip-172-31-80-126 docker]$ sudo swapoff -a
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
sudo sysctl -p
net.bridge.bridge-nf-call-iptables=1
net.bridge.bridge-nf-call-iptables = 1
```


4. Perform this ONLY on the Master machine

- Initialize kubernetes by typing below command

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all
```

```
[ec2-user@ip-172-31-80-126 docker]$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
I0913 10:32:44.629146 26680 version.go:256] remote version is much newer: v1.31.0; falling back to: stable-1.30
[init] Using Kubernetes version: v1.30.4
[preflight] Running pre-flight checks

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.80.126:6443 --token jhtgwo.4qv2vtxrcf6nvgpk \
  --discovery-token-ca-cert-hash sha256:766e48546942419274bcd18c370d2492f6e49dac9f98890804362194690f0f4a
```

- So after initialization you will get token at the end for joining master and worker. Like here I got this :(save this token as it is required later. Then you can join any number of worker nodes by running the following on each as root.)

```
kubeadm join 172.31.80.126:6443 --token jhtgwo.4qv2vtxrcf6nvgpk\
--discovery-token-ca-cert-hash
sha256:766e48546942419274bcd18c370d2492f6e49dac9f98890804362194690f0f4a
```

- Also, Copy the mkdir and chown commands from the top and execute them

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
[ec2-user@ip-172-31-80-126 docker]$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Then, add a common networking plugin called flannel file as mentioned in the code.

```
kubectl apply -f
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```



```
[ec2-user@ip-172-31-80-126 docker]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
```

- Now to Check the created pod use this command `kubectl get pods`

5. Perform this ONLY on the worker machines

Paste the below command on all 2 worker machines

- `sudo yum install iproute-tc -y`
- `sudo systemctl enable kubelet`
- `sudo systemctl restart kubelet`

Now use this

```
kubeadm join 172.31.80.126:6443 --token jhtgwo.4qv2vtxrcf6nvgpk\
--discovery-token-ca-cert-hash
sha256:766e48546942419274bcd18c370d2492f6e49dac9f98890804362194690f0f4a
```

(Optional To check the status of pods executed these commands:

`Kubectl get pods -n kube-system` :gives status of all pods

`Kubectl get daemonstat -n kube-system`: gives status of pod named daemonstat

```
[ec2-user@ip-172-31-87-114 docker]$ kubectl get pods -n kube-system
NAME                                READY   STATUS    RESTARTS   AGE
coredns-55cb58b774-fx12f           1/1     Running   0           100s
coredns-55cb58b774-xn14v           1/1     Running   0           100s
etcd-ip-172-31-87-114.ec2.internal 1/1     Running   1 (2m45s ago) 75s
kube-apiserver-ip-172-31-87-114.ec2.internal 1/1     Running   1 (2m15s ago) 2m11s
kube-controller-manager-ip-172-31-87-114.ec2.internal 0/1     CrashLoopBackOff 1 (8s ago) 70s
kube-proxy-4dv8m                   1/1     Running   2 (26s ago) 100s
kube-scheduler-ip-172-31-87-114.ec2.internal 1/1     Running   1 (2m45s ago) 76s
[ec2-user@ip-172-31-87-114 docker]$ kubectl get daemonset -n kube-system
NAME    DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
kube-proxy 1          1         1        1            1           kubernetes.io/os=linux 3m
```

)

Now to see whether master and workers get connected successfully or not run **kubectl get nodes** command on master machine

```
[ec2-user@ip-172-31-87-114 docker]$ kubectl get nodes
NAME                                STATUS   ROLES    AGE   VERSION
ip-172-31-87-114.ec2.internal      Ready    control-plane 3m21s v1.30.5
```

Conclusion: In these EC-2 instance created successfully on AWS Linux. Then I installed docker ,kuberneted and then kubelet ,kubeadm, kubectl.Then on Master machine ,I initailized the kubernetes which given me the token which will be used for connection of master and workers.then on slave I installed iproute and enabled and restarted kubelet then i enter the token which i got from master but there was an issue in joint.that is why on output i just got of only one pc mater on performing command `kubectl get nodes`.